

Virtual memory

Tran, Van Hoai

Faculty of Computer Science & Engineering
HCMC University of Technology

E-mail: hoai@hcmut.edu.vn
(*partly based on slides of Le Thanh Van*)

Outline

- 1 Background
- 2 Operations on virtual memory
- 3 Page-replacement algorithms
- 4 Frame allocation
- 5 Thrashing
- 6 Other consideration

Outline

- 1 Background
- 2 Operations on virtual memory
- 3 Page-replacement algorithms
- 4 Frame allocation
- 5 Thrashing
- 6 Other consideration

Role of virtual memory

There needs a mechanism to load **partially** a program into physical memory for execution

Why do we need to load those into memory ?

- Code to handle unusual error conditions
- Array allocated largely, but seldomly used fully

Partially loaded programs may provide some advantages

- Programmers can write program for extremely large **virtual** space
- User program takes less physical memory \Rightarrow multiprogramming increases without any sacrifice of response/turnaround time
- Less I/O needed to load/swap programs into memory

Role of virtual memory


There needs a mechanism to load **partially** a program into physical memory for execution

Why do we need to load those into memory ?

- Code to handle unusual error conditions
- Array allocated largely, but seldomly used fully

Partially loaded programs may provide some advantages

- Programmers can write program for extremely large **virtual** space
- User program takes less physical memory \Rightarrow multiprogramming increases without any sacrifice of response/turnaround time
- Less I/O needed to load/swap programs into memory




Virtual memory involves the separation of logical memory as perceived by users from physical memory

Demand paging

Idea


Demand paging is a strategy to load pages **only as they are needed** (pager is not swapper)



Valid-invalid bit is used to check **memory resident** of a page

Steps in demand paging

Pure demand paging



Pure demand paging
= only bringing pages
into memory as
required

Hardware support

- Page table: marking valid-invalid of a page
- Secondary memory (disk): known as **swap device** or **swap space**

An instruction may have page fault anywhere, restarting the **whole** instruction is needed. Consider three-address instruction **ADD A, B, C**.

- 1 Fetch and decode the instruction (ADD)
- 2 Fetch A
- 3 Fetch B
- 4 Add A and B
- 5 Store sum to C

Hardware support

- Page table: marking valid-invalid of a page
- Secondary memory (disk): known as **swap device** or **swap space**

An instruction may have page fault anywhere, restarting the **whole** instruction is needed. Consider three-address instruction **ADD A, B, C**.

- 1 Fetch and decode the instruction (ADD)
- 2 Fetch A
- 3 Fetch B
- 4 Add A and B
- 5 Store sum to C

What happens if A, B, C cannot be in main memory at the time of instruction ADD executed ?

Performance of demand paging

Denote p be probability of a page fault

$$\begin{aligned}\text{Effective access time} = & (1 - p) \times \text{memory access time} \\ & + p \times \text{page fault time}\end{aligned}$$

Page fault time generally consists of

- 1 Service the page-fault interrupt
- 2 Read in the page
- 3 Restart the process

Performance of demand paging

Denote p be probability of a page fault

$$\begin{aligned}\text{Effective access time} = & (1 - p) \times \text{memory access time} \\ & + p \times \text{page fault time}\end{aligned}$$

Page fault time generally consists of

- 1 Service the page-fault interrupt
- 2 Read in the page
- 3 Restart the process

If page fault time = 8ms and memory access time = 200ns, then

$$\text{Effective access time (in ns)} = 200 + 7,999,800 \times p$$

If **page-fault rate** $p = 1/1000$,
then effective access time = **8.2 microseconds** ($\approx 40 \times 200\text{ns}$).

Outline


- 1 Background
- 2 Operations on virtual memory**
- 3 Page-replacement algorithms
- 4 Frame allocation
- 5 Thrashing
- 6 Other consideration

Copy-on-write


Copy-on-write

- Rapid process creation
- Minimizing the number of new pages allocated to newly created processes

Before



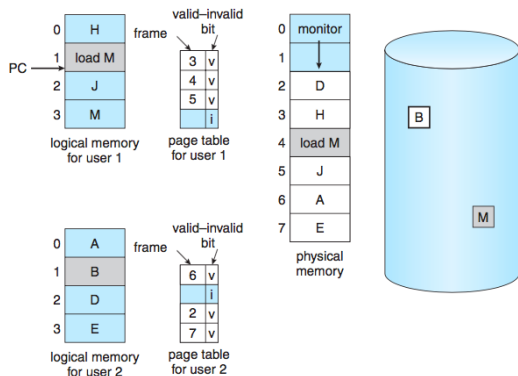
After process 1 modifies page C



- Applied for `fork()`
- `vfork()` has no copy-on-write (more efficient), and parent process is suspended

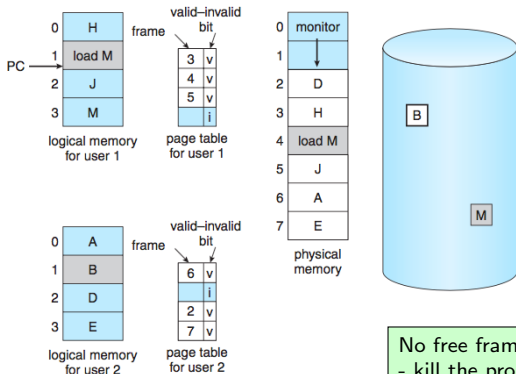
Page replacement

Some processes can utilize less pages than demanding
⇒ In order to increase multiprogramming degree, OS could **over-allocating** memory (when increasing multiprogramming degree)



Page replacement


Some processes can utilize less pages than demanding
⇒ In order to increase multiprogramming degree, OS could **over-allocating** memory (when increasing multiprogramming degree)



No free frames. Options:


- kill the process
- swap out other process
- **page replacement**

Basic page replacement



- 1 Find location of desired page on disk
- 2 Find a free frame
 - a. if there is a free frame, use it
 - b. if there is no free frame, call **page-replacement algorithm** to select a **victim frame**
 - c. Write victim to disk and change status of related data structures
- 3 Read desired page into frame, change data structures
- 4 Continue the process where page fault occurred

Basic page replacement



- 1 Find location of desired page on disk
- 2 Find a free frame
 - a. if there is a free frame, use it
 - b. if there is no free frame, call **page-replacement algorithm** to select a **victim frame**
 - c. Write victim to disk and change status of related data structures
- 3 Read desired page into frame, change data structures
- 4 Continue the process where page fault occurred

Minor tunnings

Many stuffs in above scheme can be modified to get better performance, such as

- **modify bit** (no need to write pages without any modification).
- **Read-only** pages can be discarded when desired.

Outline

- 1 Background
- 2 Operations on virtual memory
- 3 Page-replacement algorithms**
- 4 Frame allocation
- 5 Thrashing
- 6 Other consideration

Page replacement algorithm

Questions

- How many frames allocated for a process ?
 - Which frame to be replaced ?
 - Different programs have different memory references (namely **reference string**)
-
- Recording an address sequence as follows
0100, 0432, 0101, 0612, 0102, 0103, 0104, 0101, 0611, 0102, 0103, 0104,
0101, 0610, 0102, 0103, 0104, 0609, 0102, 0105
 - 100 bytes per page, then address sequence converted to reference string
1, 4, 1, 6, 1, 6, 1, 6, 1, 6, 1

FIFO page replacement

Let's try the following reference string for **three-frame** memory

7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

FIFO page replacement

Let's try the following reference string for **three-frame** memory

7, 0, 1, 2, 0, 3, 0, 4, 2, 3, 0, 3, 2, 1, 2, 0, 1, 7, 0, 1

FIFO page replacement

Choose next page to be replaced in **FIFO** scheme

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2	2	2	4	4	4	0	0	0	0	0	0	0	0	7	7	7
	0	0	0	3	3	3	2	2	2	1	1	3	3	1	1	1	1	0	0
		1	1	1	0	0	0	0	3	3	2	2	2	2	2	2	2	2	1


page frames

15 page faults

Optimal replacement algorithm

Bélády's anomaly for FIFO


Number of available frames increases, but page fault does **increase** (*proved to be unbounded (2010)*)



Optimal replacement algorithm

Bélády's anomaly for FIFO

Number of available frames increases, but page fault does **increase** (*proved to be unbounded (2010)*)



Optimal (MIN/OPT) page replacement

Replace the page that **will** not be used for the **longest** period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

Optimal replacement algorithm

Bélády's anomaly for FIFO

Number of available frames increases, but page fault does **increase** (*proved to be unbounded (2010)*)



Optimal (MIN/OPT) page replacement

Replace the page that **will** not be used for the **longest** period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

9 page faults

Optimal replacement algorithm

Bélády's anomaly for FIFO

Number of available frames increases, but page fault does **increase** (proved to be unbounded (2010))



Optimal (MIN/OPT) page replacement

Replace the page that **will** not be used for the **longest** period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



9 page faults

There is **no information** of the future

LRU replacement algorithm

LRU page replacement

LRU (Least recently used) chooses the page that **has not been used** for the longest period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1		
	0	0	0		0		0	0	3	3			3		0		0		
		1	1		3		3	2	2	2			2		2		7		

page frames


LRU replacement algorithm

LRU page replacement

LRU (Least recently used) chooses the page that **has not been used** for the longest period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1




page frames

LRU replacement algorithm

LRU page replacement

LRU (Least recently used) chooses the page that **has not been used** for the longest period of time



We need hardware support to implement LRU algorithm.

- **Counters:** an additional field in each page-table entry to store **time-of-use** (clock or counter). Referencing a page → copying storing clock register to time-of-use field
- **Stack:** using the idea of a stack, **"Last In First Out"**

Other algorithms

- LRU-approximation page replacement
 - Additional-reference-bits algorithm
 - Second-chance algorithm
 - Enhanced second-chance algorithm
- Counting-based page replacement
 - Least frequently used (LFU)
 - Most frequently used (MFU)
- Page-buffering algorithm

Outline

- 1 Background
- 2 Operations on virtual memory
- 3 Page-replacement algorithms
- 4 Frame allocation**
- 5 Thrashing
- 6 Other consideration

Minimum number of frames

Minimum number of frames is defined by **computer architecture**

IBM MVC instruction needs

- a) 6 bytes (can straddle 2 pages)
- b) source location (2 pages)
- c) destination location (2 pages)

→ 6 pages totally

Number levels of indirection should be limited

Allocation algorithms

There are m frames

- **Equal allocation**: n processes $\rightarrow \lfloor \frac{m}{n} \rfloor$ for each process
- **Proportional allocation**: virtual memory size of p_i is s_i

$$a_i = m \times s_i / \sum s_i$$

a_i possibly depends on a combination of size and **priority** of the process

Allocation algorithms

There are m frames

- **Equal allocation**: n processes $\rightarrow \lfloor \frac{m}{n} \rfloor$ for each process
- **Proportional allocation**: virtual memory size of p_i is s_i

$$a_i = m \times s_i / \sum s_i$$

a_i possibly depends on a combination of size and **priority** of the process

Performance issues

- Global vs. local allocation
 - Global replacement: replacement frame is any in the set of **all** frames
 - Local replacement: replacement frame is in its **own set of allocated** frames
- Non-uniform memory access (NUMA)

Outline

- 1 Background
- 2 Operations on virtual memory
- 3 Page-replacement algorithms
- 4 Frame allocation
- 5 Thrashing**
- 6 Other consideration

Thrashing


A process has **no enough** frames. Page faults occur frequently.
Time for paging is more than that for execution.

Thrashing

A process has **no enough** frames. Page faults occur frequently. Time for paging is more than that for execution.

Thrashing

High paging activity is called **thrashing**




Working with thrashing

We can limit effects of thrashing by

- Using **local replacement algorithm**, not steal frames from another process
- **Working-set strategy**: using a **locality model** (locality = set of pages that are actively **used together**)

A function has its own locality which consists of

- its local variables
- subset of global variables





Working-set model

- Δ : working-set window
- Working set: set of pages in the most recent Δ page references

$$\Delta = 10$$

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...




Denote WSS_i be working-set size of process i . Total demand is

$$D = \sum WSS_i$$

We can monitor D and respond accordingly when $D > m$

Page-fault frequency (PFF)



Outline

- 1 Background
- 2 Operations on virtual memory
- 3 Page-replacement algorithms
- 4 Frame allocation
- 5 Thrashing
- 6 Other consideration**

- **Prepaging:** prevent large number of page faults that occur **when a process is started**
- **Page size:** choice of page size is often performed in OS design
- **TLB reach**
- **Inverted page tables**
- **Program structure**

```
int i, j;  
int data[128][128];  
  
for( j = 0; j < 128; j++ )  
    for( i = 0; i < 128; i++ )  
        data[i][j] = 0;
```

```
int i, j;  
int data[128][128];  
  
for( i = 0; i < 128; i++ )  
    for( j = 0; j < 128; j++ )  
        data[i][j] = 0;
```

- **Prepaging:** prevent large number of page faults that occur **when a process is started**
- **Page size:** choice of page size is often performed in OS design
- **TLB reach**
- **Inverted page tables**
- **Program structure**

```
int i, j;  
int data[128][128];  
  
for( j = 0; j < 128; j++ )  
    for( i = 0; i < 128; i++ )  
        data[i][j] = 0;
```

```
int i, j;  
int data[128][128];  
  
for( i = 0; i < 128; i++ )  
    for( j = 0; j < 128; j++ )  
        data[i][j] = 0;
```

If OS allocates **fewer than 128 frames of 128 words** for entire program, the left code has **$128 \times 128 = 16,384$ page faults**. The right code possibly has only **128 page faults**.